

Lecture 10. Transaction

What is a Transaction?

Transaction = a **single logical unit of work** consisting of one or more SQL statements.

- All operations inside a transaction are treated as one package:
- Either **all succeed** → changes are saved
- Or **any fails** → all changes are undone

PostgreSQL transactions follow ACID properties (Atomicity, Consistency, Isolation, Durability)

ACID Properties

Atomocity

"All or nothing"

If **any operations** in the transaction fails:

- Whole transaction is **rolled back**
- **No partial changes** remain in the database

Consistency

Transaction moves database from **one valid state to another**

All constraint must hold:

- primary/unique keys
- foreign keys, referential integrity
- CHECK constraint, business rules

Example: likes count in a blog must always match actual rows in "likes" table

Isolation

- Concurrent transactions must **not interfere** with each other.
- A running transaction behaves **as if it is alone in the system**, from its point of view.
- Intermediate (uncommitted) changes of one transaction are **not visible** to others.

Durability

Once a transaction is **committed**, its changes are:

- **Permanent** (written to non-volatile storage / WAL).
 - Survive crashes, power loss, restarts.
-

Transaction Lifecycle

Phases:

1. **BEGIN** - start transaction
2. **EXECUTE** - run SQL queries (SELECT/INSERT/UPDATE/DELETE, etc.).
3. **COMMIT** - save all changes if everything OK
4. **ROLLBACK** - undo all changes if error / decisions to cancel

PostgreSQL treats **every SQL statement** as running in a transaction:

- Without explicit `BEGIN`, each statement runs in its own implicit transaction
(`BEGIN` + `COMMIT`)
-

Basic Transaction Control Statement

BEGIN

Start a transaction:

```
BEGIN;  
--or  
BEGIN TRANSACTION;  
--or  
BEGIN WORK;
```

Everything after `BEGIN` until `COMMIT` / `ROLLBACK` is **one transaction block**.

COMMIT

Apply (save) changes:

```
COMMIT;  
--or  
COMMIT WORK;  
--or  
COMMIT TRANSACTIONS;
```

After `COMMIT` :

- Changes are visible to other sessions
- Durability guarantee kicks in

ROLLBACK

Undo all changes made since `BEGIN` :

```
ROLLBACK;  
--or  
ROLLBACK WORK;  
--or  
ROLLBACK TRANSACTION;
```

Use when:

- An error occurs.
- You decide not to apply the changes.

Isolation Levels

PostgreSQL supports standard **transaction isolation levels**:

- **Read Uncommitted**
 - Can *theoretically* read uncommitted data.
 - In PostgreSQL, this behaves like **Read Committed** (dirty reads are still not allowed).
- **Read Committed** (Default)
 - Each query sees **only committed data** at the moment that query starts.
 - If another transaction commits in between your queries, your next query can see its changes.
- **Repeatable Read**
 - Transaction sees a **snapshot** of the database taken at the start of the transaction.
 - The same row read twice in the same transaction returns **the same data**, even if others commit updates.
- **Serializable**
 - Strictest level.
 - Guarantees transactions behave **as if executed one by one in some serial order**.
 - Prevents subtle anomalies at the cost of more blocking or rollbacks.

Concurrency Anomalies

1. Dirty Read

- Transaction T1 reads data written by T2 **that is not committed yet**.
- If T2 rolls back, T1 has seen “ghost” data.
- PostgreSQL avoids this even at Read Uncommitted.

2. Non-Repeatable Read

- T1 reads a row.
- T2 commits an update to that row.
- T1 reads the same row again → gets **different values**.

3. Phantom Read

- T1 runs a query with some WHERE condition (e.g. `salary > 1000`) and gets N rows.
- T2 inserts/deletes rows matching that condition and commits.
- T1 runs the **same query** again → **different set of rows** (new “phantoms” appear/disappear).

4. Lost Update

- T1 and T2 both read the same value, then both write based on that old value.
- One update overwrites the other → one update is “lost”.

5. Write Skew

- T1 and T2 read overlapping data and both decide to update **different** rows based on a shared condition.
- Individually each transaction is “valid”, but combined they violate a business rule.

Setting Isolation Level

Global per session:

```
SET TRANSACTION ISOLATION LEVEL SERIALIZABLE;
BEGIN;
--operations
COMMIT;
```

Or:

```
BEGIN TRANSACTION ISOLATION LEVEL REPEATABLE READ;
--operations
```

```
COMMIT;
```

Isolation level affects **visibility** and **allowed anomalies**.

Savepoints

| **Savepoint** = checkpoint **inside** a transaction

Use when you want **partial rollback**

Creating a savepoint:

```
BEGIN;
```

```
-- some operations
```

```
SAVEPOINT sp1;
```

```
-- more operations
```

```
ROLLBACK TO sp1; -- undo only changes after sp1
```

```
COMMIT; -- commit everything before sp1 + after rollback fixes
```

Key points:

- `SAVEPOINT name;` marks a position.
- `ROLLBACK TO name;` :
 - discards changes from after `SAVEPOINT` to the rollback point,
 - **keeps earlier work in the same transaction.**
- After `ROLLBACK TO`, the savepoint still exists and can be reused.
- `RELEASE SAVEPOINT name;` :
 - removes the savepoint; cannot roll back to it anymore.
 - Also automatically releases all savepoints defined after it.

Benefits:

- More flexible error handling (rollback only part of transaction).
 - Better modularity (complex transaction broken into steps).
 - Useful for financial/booking systems where some steps may fail but you don't want to discard everything.
-

When to Use Transactions

- **Financial applications** (banking, e-commerce payments).
 - **Inventory management** (adjusting stock, orders).
 - **Reservation systems** (airlines, hotels).
 - Any **multi-user application** updating shared data.
-

Advantages and Disadvantages

Advantages

- **ACID guarantees** → data integrity and consistency.
- **Error handling via ROLLBACK**.
- **Improved concurrency control** with proper isolation.
- Combine multiple operations into one unit → can also be more efficient.

Disadvantages

- **Complexity & overhead:**
 - Application logic more complex.
 - Extra work for the DB engine (locking, logging).
- **Reduced concurrency** (depending on isolation level and locks).
- **More resource usage** for long or complex transactions.

- **Harder debugging** (many operations and dependencies inside one transaction).