

# Lecture 11. Stored Procedures & PL/pgSQL

## What is a Stored Procedure?

A **stored procedure** is a **named block of SQL + procedural logic** stored in the database and executed on the server.

It can contain:

- Multiple SQL statements
- Variables
- Loops, conditions
- Error handling

Main goals: **encapsulate logic, reuse code, improve performance, centralize business rules.**

---

## Advantages of Stored Procedures

- **Fewer round trips**
  - Instead of sending many separate SQL statements from the application to the DB, the app calls **one procedure**.
  - DB executes all internal SQL on the server side.
- **Better performance**
  - User-defined procedures are **precompiled / cached** on the server.
  - Execution plan can be reused → less overhead.
- **Reuse across applications**
  - Once written, the procedure can be called from **any app** that connects to that database

# Disadvantages of Stored Procedures

- **Slower development:**
    - Requires DB-specific skills (PL/pgSQL, SQL inside DB).
  - **Harder versioning & debugging:**
    - Logic sits inside database, not in normal application codebase.
  - **Portability issues:**
    - Procedures are DB-specific (PostgreSQL vs MySQL vs SQL Server).
    - Migration to another DBMS may require rewriting.
- 

## Procedural Languages in PostgreSQL

PostgreSQL supports several languages for writing stored procedures / functions:

- Built-in:
  - **SQL**
  - **PL/pgSQL**
  - **C**
- Via extensions:
  - **PL/Perl, PL/Python, PL/JavaScript, PL/Ruby, PL/Java, PL/Tcl**, etc.

In this lecture, focus is on **PL/pgSQL**.

---

## PL/pgSQL Overview

**PL/pgSQL (Procedural Language / PostgreSQL)** is PostgreSQL's main procedural language.

It allows:

- Writing **functions, procedures, triggers**

- Using control-flow logic:
  - `IF`, `CASE`, `LOOP`, `WHILE`, `FOR`
- Using **variables, parameters, error handling**
- Executing **dynamic SQL** at runtime
- Deep integration with SQL (easy to call queries inside code)

## Key Features:

- **Control Structures:** `IF`, `CASE`, `LOOP`, `WHILE`, `FOR`
  - **Error Handling:** `EXCEPTION` blocks
  - **Dynamic SQL:** `EXECUTE`
  - **Variables and Parameters:** `DECLARE` block
  - **Integration:** runs inside PostgreSQL server, close to the data.
- 

## Basic PL/pgSQL Block Structure

Generic shape:

```

CREATE OR REPLACE PROCEDURE proc_name(arg_list)
LANGUAGE plpgsql
AS $$

DECLARE
    -- variable declarations
    total_sales numeric := 0;

BEGIN
    -- main code
    -- SQL statements, IF, loops, etc.

    -- optional error handling

EXCEPTION
    WHEN division_by_zero THEN
        RAISE NOTICE 'Error: division by zero';

```

```
END;  
$$;
```

Important elements:

- `LANGUAGE plpgsql` – tells PostgreSQL to use PL/pgSQL.
- `$$ ... $$` – delimiters for the procedure body.
- `DECLARE` – section where you declare variables (optional).
- `BEGIN ... END` – main block.
- `EXCEPTION ...` – handles runtime errors (optional).

---

## Variables (DECLARE Section)

In the `DECLARE` block you define local variables:

```
DECLARE  
    counter    integer := 0;  
    customer_id integer;  
    total_amount numeric;
```

- You can set default values with `:=`.
- Variables can be used in SQL statements or expressions inside the procedure.

---

## Control Structure (Logic):

PL/pgSQL supports classic control-flow:

- `IF ... THEN ... ELSE ... END IF;`
- `CASE ... WHEN ... THEN ... END CASE;`
- Loops:
  - `LOOP ... END LOOP;`

- WHILE condition LOOP ... END LOOP;
- FOR var IN 1..10 LOOP ... END LOOP;

Used to implement complex business rules directly in the DB.

---

## Error Handling with EXCEPTION

EXCEPTION blocks allow you to **catch and handle errors**:

```
BEGIN
    -- risky operations
    UPDATE accounts
    SET balance = balance - 100
    WHERE id = 1;

EXCEPTION
    WHEN foreign_keyViolation THEN
        RAISE NOTICE 'Error: foreign key violation';
END;
```

- Prevents the whole transaction from failing uncontrolled.
  - You can log errors, revert part of logic, or raise custom messages.
- 

## Parameters in Stored Procedures

Parameters allow passing values **into** and **out of** a procedure.

PostgreSQL supports:

### 1. IN

- Default type if not specified (for procedures).
- Passes a value **into** the procedure.
- Cannot be modified to return something back.

### 2. OUT

- Used to **return values** from the procedure.
- Act like extra result columns.

### 3. INOUT

- The parameter is passed in and can be **modified**.
- Final value is returned to the caller.

Conceptually:

```
CREATE PROCEDURE adjust_salary(IN emp_id int, IN delta numeric)
```

...

```
CREATE PROCEDURE get_salary(IN emp_id int, OUT salary numeric)
```

...

```
CREATE PROCEDURE change_and_report(INOUT amount numeric)
```

...

## Function vs Stored Procedures

### Functions

- **Must return** a value (scalar, row, table, or `void`).
- Typically used in:
  - `SELECT`, `WHERE`, `JOIN`, etc.
- **No transaction control** inside:
  - Cannot `COMMIT` or `ROLLBACK` within a function.

### Stored Procedures

- May **return no value**.
- Invoked with `CALL procedure_name(...)`.
- **Support transaction control** inside:
  - Can use `COMMIT`, `ROLLBACK`, `SAVEPOINT` etc. in certain patterns.

- Better suited for **multi-step workflows** that need control over transactions.
- 

## Dynamic SQL in Stored Procedures

| **Dynamic SQL** = constructing SQL text at runtime and executing it.

Usage pattern:

```
EXECUTE 'DELETE FROM ' || quote_ident(table_name) ||
        ' WHERE created_at < $1'
    USING cutoff_date;
```

- Allows building queries where table name, columns, or conditions are dynamic.
- Very powerful but must be used carefully (SQL injection, complexity).